

KOVVURI RISHI KESAV REDDY

Hyderabad, India | +91-9030124944 | rishikesav2005@gmail.com | [LinkedIn](#) | [GitHub](#) | [LeetCode](#)

PROFESSIONAL SUMMARY

B.E. student in Artificial Intelligence & Data Science at CBIT (CGPA: 9.45/10, expected graduation 2027) with hands-on experience building and deploying full-stack applications and ML systems. Designed a stacked gradient boosting ensemble (LightGBM, XGBoost, CatBoost) for product price prediction in Amazon ML Challenge 2025, a national-level competition. Deployed ZYVA, a live healthcare platform with a RESTful API backend, OAuth2/JWT authentication, role-based access, and cloud hosting, serving real users. Looking to apply software engineering and applied ML skills to large-scale, production systems.

EDUCATION

Chaitanya Bharathi Institute of Technology (CBIT)

B.E., Artificial Intelligence & Data Science | Gandipet, Hyderabad

2023 – Present

CGPA: 9.45 / 10 (Semester 5)

Sri Chaitanya Junior College | Intermediate (MPC) | 2021-2023

975 / 1000 (97.5%)

Prachin Global School | CBSE (10th) | 2020-2021

462/500 (92.4%)

TECHNICAL SKILLS

Languages: Python, SQL, HTML

ML / DL: Scikit-learn, TensorFlow, Pandas, Numpy, Matplotlib, Feature Engineering, K-Fold CV, Ensemble Methods

Web & Backend: RESTful APIs, Google OAuth2, JWT, Node.js, Express.js, FastAPI

Cloud & Deployment: Render (cloud hosting), Oracle Cloud Infrastructure (AI Foundations Certified)

Databases: MongoDB (GridFS), MySQL

Tools: Git, GitHub, Jupyter Notebook, VS Code, Kaggle

Core CS: Data Structures & Algorithms, OOPs, Problem Solving, Cloud Computing

EXPERIENCE

Machine Learning Participant — Amazon ML Challenge 2025

2025

- Designed and optimized regression models on a noisy, large-scale real-world product dataset under competitive time and evaluation constraints.
- Built a stacked ensemble (LightGBM, XGBoost, CatBoost) with weighted blending to improve generalization over single-model baselines.
- Engineered price-quantile stratified cross-validation to handle skewed price distributions and reduce overfitting.

External Relations Head – AgentForce Club, CBIT

2025-Present

- Presented partnership proposals and coordinated with cross-functional teams to execute events successfully.
- Increased event participation by 40% through targeted outreach campaigns and planning the event according to the target audience.
- Collaborated with cross-functional teams to ensure seamless execution of club initiatives.

PROJECTS

Mentora AI

2026

Python, FastAPI, Groq LLM, Sentence-transformers, ChromaDB, RAG, librosa, PyMuPDF, python-docx AI-Powered Full Stack Application

- Designed and built a GenAI-powered RAG tutoring system that ingests student notes (PDF/DOCX) into ChromaDB using sentence-transformer embeddings and retrieves semantically relevant context for Groq LLM, with a 0.30 cosine similarity confidence gate to reduce hallucinations.
- Implemented a hybrid NLP-based mood classifier combining semantic similarity, keyword matching, and behavioural signals (response latency, backspace count, attempt number) to classify five student emotional states without labelled training data.
- Engineered a 4-dimensional tone space (frustration level, confidence boost, pace slowdown, encouragement intensity) that dynamically constructs system prompts to personalize LLM teaching behaviour for each detected mood.
- Built a knowledge gap tracker with spaced repetition that records chat interactions and quiz grades, normalizes topics, identifies weak concepts, and schedules revision using the SM-2 algorithm.
- GitHub: <https://github.com/KovvuriRishiKesavReddy/MentoraAI>

- Built a pipeline merging structured (ATS JSON) and unstructured (free-text) candidate records into one canonical, schema-validated profile per candidate, with full field-level provenance and confidence scoring.
- Designed a source-priority conflict-resolution policy and a multi-tier candidate-matching strategy (exact email/phone → fuzzy name fallback) to merge records across sources without data loss or silent corruption.
- Implemented a runtime-configurable projection layer decoupled from the core merge engine, allowing the same pipeline to emit different output schemas with zero code changes.
- Covered extraction, normalization, merge logic, and projection with 21 unit tests against deliberately engineered edge cases (conflicting sources, missing fields, malformed input).
- GitHub: <https://github.com/KovvuriRishiKesavReddy/Multi-Source-Candidate-Data-Transformer>

- Designed and deployed a full-stack healthcare platform with secure authentication, appointment management, and prescription storage. Processed and managed healthcare-related data securely to support appointment scheduling, prescription management, and user workflows.
- Used GridFS for medical file storage and role-based access control for patients and doctors; deployed on Render for public accessibility.
- Designed and implemented a RESTful API architecture with 25+ endpoints supporting real-time booking, authentication, and resource management.
- Live: zyva-healthcare-utus.onrender.com

- Trained GPU-accelerated LightGBM, XGBoost, and CatBoost across 7 stratified folds with price-quantile binning.
- Optimized ensemble blend weights using SciPy SLSQP on out-of-fold predictions, with Lasso meta-stacking for the final layer.
- Evaluated models using SMAPE and optimized performance through systematic experimentation.
- Link: <https://www.kaggle.com/code/krkredd/amazon-ml>

CERTIFICATIONS

- **Honeywell Future Skills Edge Program on AI** | Honeywell (2026)
- **Oracle Cloud Infrastructure — AI Foundations Associate** | Oracle (2025)
- **Associate Developer Certification** | MongoDB University (2025)
- **Web Technology Certification** | Internshala (2024)

ACHIEVEMENTS

- **Panel Winner, Hacktoberfest CBIT** — Recognised by judging panel at CBIT's annual open-source hackathon.
- **Solved 350+ algorithmic and data structure problems on LeetCode.**
- **93% accuracy on the Clothing Size Recommendation System ensemble.**